

Intent-Driven Spatial Search Engine

A Hybrid NL-to-SQL Architecture for Conversational Geographic Queries

BITS Pilani — Visiting Summer Research Internship (VISRI) Project

July 2, 2026

Contents

1	Introduction	2
1.1	Background	2
1.2	Motivation	2
1.3	Objective	2
2	Related Work (with Comparisons)	3
2.1	Traditional Geocoders and Commercial Maps	3
2.2	AI Recommendation Systems	3
2.3	The Challenge with Standard Text-to-SQL AI	3
2.4	Our Proposed Solution: The Hybrid Architecture	4
3	Our Evaluation Framework	4
3.1	Evolution of Text-to-SQL Evaluation in Literature	4
3.2	Step 1: Code Structure Check	5
3.3	Step 2: Live Database Execution	5
4	Dataset Curation & Methodology	6
4.1	Database Construction (OpenStreetMap)	6
4.2	Justification for Zero-Shot Benchmark Scale (Diversity Over Volume)	6
4.3	The English SQL Benchmark (50 Queries)	6
4.4	The Multi-Lingual Translation Dataset (56 Queries)	7
5	Evaluation and Result(s) (Comparisons)	7
5.1	Text-to-SQL Model Selection	7
5.1.1	Conclusion and Root Cause Analysis	7
5.2	Testing the Translation AI	8
5.2.1	Why Sarvam Won	8
5.3	Testing Voice Recognition (ASR)	8
5.3.1	Why SeamlessM4T Won for Voice	8
5.4	The Final End-to-End App	9
6	Conclusion	9
6.1	Future Work	10

Abstract

Traditional digital maps rely on rigid keyword searches, which often fail when users ask for things conversationally. To solve this, we built an Intent-Driven Spatial Search Engine that translates colloquial, multi-lingual queries directly into structured PostGIS database commands. Running 100% offline on a local GPU, the system uses SeamlessM4T for speech recognition and translation, paired with Qwen2.5-Coder to act as a Text-to-SQL engine and intent router. We rigorously evaluated multiple AI models using our custom Hybrid Evaluation Framework, which validates both the SQL structure (using Abstract Syntax Trees) and the live database results. Our final architecture achieved a 78.0% success rate on complex geographic intents, proving it can successfully bridge the gap between abstract human intent (like “a quiet place to study”) and rigid geographic databases.

1 Introduction

1.1 Background

Digital mapping systems power everything from ride-sharing to daily navigation. However, despite advancements in satellite imagery, their search engines are still surprisingly limited. Mainstream mapping apps (like Google Maps or Apple Maps) rely heavily on exact keyword filtering.

While these apps are great at finding exact addresses or formal business names, they fail when users search for descriptive, intent-driven locations. People naturally think about spaces based on their needs—such as “a quiet place to study” or “a late-night pharmacy.” When forced through a traditional map search, these conversational queries break down.

1.2 Motivation

This limitation causes what we call a **Local Intent Failure**. For instance, if a student on campus searches for “quiet places near me,” traditional systems do not understand that “quiet” means a library or a park. Instead, they look for a business literally named “Quiet.” This leads to the map suggesting a “Quiet Place” restaurant located 1,000 kilometers away, completely ignoring the “near me” spatial constraint.

This problem is even worse in India. Indian spatial navigation is highly descriptive, informal, and frequently code-mixed (e.g., Hinglish). A colloquial query like “koi chai ki tapri hai kya paas mein?” (Is there a tea stall nearby?) requires an engine that can translate the slang, identify the location category, and strictly enforce a local geographic radius. Traditional mapping engines fail on all three fronts.

1.3 Objective

To solve this, we need a fundamental shift in how spatial search works. The objective of this project is to abandon rigid keyword-matching entirely. Instead, we use Large Language Models (LLMs) to translate human intent directly into mathematical spatial queries executed by a relational database.

Specifically, this project aims to:

1. **Dynamic Intent Capturing:** Utilize specialized Text-to-SQL models to translate conversational queries into mathematically enforced spatial constraints (PostGIS SQL), ensuring the database executes a bounding box rather than relying on a traditional text-search engine to guess what “near me” means.
2. **Colloquial Multilingual Translation:** Integrate state-of-the-art open-weight translation models (like Sarvam-Translate) to bridge the gap between regional Indian dialects/slang and the formal English syntax required by the SQL generation AI.

3. **Empirical Evaluation:** Develop a robust Hybrid Evaluation Framework that proves the accuracy of LLM-generated spatial SQL against messy, real-world data, scoring for both structural soundness and semantic accuracy.
4. **End-to-End System Integration:** Deploy a functional, 100% offline pipeline with an intuitive conversational User Interface (Gradio + Folium), allowing users to speak naturally (via SeamlessM4T) and view their results dynamically rendered on a map alongside locally synthesized native-language audio responses.

By achieving these objectives, the system will accurately map human-centric vocabulary to underlying geographic realities, restoring the integrity of core functionalities like spatial proximity.

2 Related Work (with Comparisons)

The development of intuitive spatial search engines has gained significant attention, though approaches diverge heavily between industry standards and academic research.

2.1 Traditional Geocoders and Commercial Maps

Current mapping engines rely on exact keyword matching and structured address lookups.

- **Traditional Geocoders (e.g., Nominatim):** The default search engine for OpenStreetMap requires perfectly structured addresses (Country, State, City, Street). It is completely blind to conversational language, meaning it fails immediately if a user searches for a description rather than an exact address.
- **Commercial Maps (e.g., Google Maps):** These platforms use standard text-ranking algorithms. They suffer from **Exact-Match Bias**. Because they treat queries as simple text strings rather than understanding the user’s actual intent (like needing a quiet place nearby), they frequently suggest highly irrelevant, distant locations just because the business name matches the query.

2.2 AI Recommendation Systems

To fix the exact-match problem, researchers have built AI recommendation systems that try to capture the “vibe” of a location using text embeddings.

- **Graph-Based Recommenders:** While these systems are good at understanding descriptive intent, they suffer from **Weak Spatial Enforcement**. Because they are text-based AI models, they treat distance as just another text feature. They cannot execute strict mathematical boundaries, meaning they struggle to enforce a hard 3-kilometer radius around the user.

2.3 The Challenge with Standard Text-to-SQL AI

Modern AI models are excellent at translating English into standard database SQL (e.g., for finance or logistics). However, they frequently fail when asked to write **GeoSQL** (like PostGIS). Standard models struggle to understand the complex mathematical syntax required for geographic routing, such as calculating bounding boxes or radii (**ST_DWithin**). Without heavy fine-tuning, most standard AI coders are too unreliable for map routing.

2.4 Our Proposed Solution: The Hybrid Architecture

This project solves the problems of both commercial maps and text-based AI recommenders. We use AI not to search text, but to write structured SQL that runs directly on a geographic database (PostGIS).

Instead of doing a text search for a business named “Quiet,” our AI translates “quiet” into specific location tags (like a library) and translates “near me” into a strict mathematical formula (`ST_DWithin`). Because the AI generates database code rather than text guesses, the database can enforce a 100% rigid geographic boundary. This perfectly combines the deep understanding of AI with the mathematical precision of a traditional database.

3 Our Evaluation Framework

Testing an AI’s ability to write SQL is difficult. You cannot simply check if the AI’s code exactly matches an answer key, because SQL is highly flexible. Two completely different queries might return the exact same correct data.

To accurately test our models against messy, real-world data, we built a custom two-step **Hybrid Evaluation Framework**. The design of this framework was directly informed by the evolution of evaluation methodologies in the Text-to-SQL literature, which we summarize below.

3.1 Evolution of Text-to-SQL Evaluation in Literature

The academic community has iterated through several generations of evaluation strategies, each addressing the limitations of the previous one.

1. **Exact Match (WikiSQL, 2017):** The earliest large-scale benchmark, WikiSQL [1], contained 87,726 question-SQL pairs but used only simple, single-table queries. Its primary metric was *Exact Match (EM)*, which checks if the predicted SQL string is character-for-character identical to the gold standard. This was widely criticized because multiple, syntactically different SQL queries can return the same correct result.
2. **Component Matching (Spider, 2018):** The Spider benchmark [2] introduced complex, cross-domain queries (10,181 examples across 200 databases) with multi-table JOINS, nested subqueries, and GROUP BY clauses. To overcome the rigidity of Exact Match, Spider introduced *Exact Set Match*, which parses the SQL into individual clauses (SELECT, WHERE, ORDER BY) and compares the predicted *sets* of columns, tables, and operators against the gold standard. **Our structural scoring (Step 1) is directly inspired by this approach**, using the `sqlglot` library to parse SQL into Abstract Syntax Trees (ASTs) for clause-level comparison.
3. **Execution Accuracy (BIRD, 2023):** The BIRD benchmark [3] (12,751 pairs across 95 real-world databases totaling 33.4 GB) introduced *Execution Accuracy (EX)*, which validates queries by running them against a live database and comparing the returned result sets. BIRD also introduced the *Valid Efficiency Score (VES)* to reward computationally efficient queries. **Our live execution scoring (Step 2) follows this execution-based philosophy**, but we extend it from a binary pass/fail to a continuous *Semantic F1 Score* that rewards partial spatial reasoning.
4. **Robustness Testing (Dr.Spider, 2023):** Dr.Spider [4] demonstrated that high-performing models are fragile, suffering up to 50.7% accuracy drops when subjected to 17 types of linguistic and schema perturbations. **Our multilingual benchmark follows this perturbation philosophy**: each of the 7 base spatial intents is tested across 8 linguistic variants

(native script, Romanized, code-mixed slang across 6 languages), effectively functioning as a robustness perturbation suite.

5. **Spatial SQL Evaluation (GeoSQL-Eval, 2026):** Most recently, GeoSQL-Eval [5] became the first benchmark specifically designed for PostGIS spatial queries, containing 14,178 tasks across 340 spatial functions and 82 databases. It found that approximately 70% of LLM errors in spatial tasks stem from incorrect spatial function selection (e.g., hallucinating `ST_Buffer` instead of `ST_DWithin`), rather than basic SQL syntax errors. Our project addresses the same spatial reasoning challenge within a single, dense OpenStreetMap schema.

Our Hybrid Evaluation Framework synthesizes the strengths of these approaches: *clause-level structural parsing* from Spider, *live database execution* from BIRD, and *multilingual perturbation testing* from Dr.Spider, applied to the *spatial SQL domain* identified by GeoSQL-Eval.

3.2 Step 1: Code Structure Check

Before executing the generated SQL on the live database (which could result in crashes or massive unbounded data downloads), we first perform a strict Abstract Syntax Tree (AST) structural analysis. Using the `sqlglot` parser, we decompose the SQL query into its constituent clauses and evaluate it across seven independent dimensions:

- **SELECT Clause Accuracy:** Verifies that the correct geographic and attribute columns are projected without hallucinating non-existent columns.
- **FROM Clause Mapping:** Ensures the query targets the correct base table (`campus_places`) and uses appropriate aliases for self-joins.
- **WHERE Clause Filtering:** Confirms that mandatory filters (e.g., `name IS NOT NULL`) and required attribute tags are present to prevent unbounded data retrieval.
- **WHERE_OPS Logical Safety:** Checks the correct usage of logical operators, ensuring `ILIKE` is used for loose text matching and `AND/OR` groupings are logically sound.
- **ORDER BY Accuracy:** Validates that results are sorted correctly, which is especially critical for nearest-neighbor spatial queries sorted by distance.
- **LIMIT Constraints:** Ensures that the query strictly adheres to pagination rules (e.g., `LIMIT 1` or `LIMIT 5`) to prevent flooding the UI.
- **Spatial Function Accuracy:** Inspired by the GeoSQL-Eval framework (2026), we verify that the model selects the correct PostGIS mathematical function (e.g., `ST_DWithin` for radii vs `ST_Distance` for nearest-neighbor) and correctly casts geometries to geographies (`::geography`) for the specific spatial intent.

3.3 Step 2: Live Database Execution

If the code passes the safety check, we execute it live on the PostGIS database. We then compare the locations the AI found against the correct “Gold Standard” locations from our benchmark. We score the AI based on two metrics:

- **Precision:** Out of all the locations the AI returned, how many were actually right?
- **Recall:** Out of all the correct locations that exist, how many did the AI successfully find?

The Passing Grade (50% Rule): Because real-world map data is inherently messy and inconsistent, requiring the AI to achieve a perfect 100% match is unrealistic. For example, if a user asks for “sports facilities,” and the AI successfully finds 9 out of 10 correct locations but misses one due to a minor syntax choice, it is still a highly successful search. Therefore, our framework dictates that any AI query that correctly identifies at least half of the target locations ($\text{Score} \geq 0.50$) receives a **PASS**. Anything below is a **FAIL**.

If the AI writes broken code that crashes the database, it immediately receives a **FAIL**, ensuring strict accountability.

4 Dataset Curation & Methodology

To rigorously test both the spatial reasoning of the Text-to-SQL models and the linguistic robustness of the translation models, two highly specialized datasets were manually curated for the Pilani and IIT Bombay campus regions.

4.1 Database Construction (OpenStreetMap)

We needed a live, real-world map to test our AI against. We pulled mapping data directly from OpenStreetMap for a 3km radius around the BITS Pilani and IIT Bombay campuses. We filtered for relevant map features (like cafes, parks, and libraries), cleaned the data, and loaded it into a local PostgreSQL spatial database. This became the playground for our AI models.

4.2 Justification for Zero-Shot Benchmark Scale (Diversity Over Volume)

Because we evaluate frontier LLMs (like Qwen2.5-Coder) in a Zero-Shot setting, inflating the test set with hundreds of queries would inevitably introduce structural redundancy (asking the same SQL logic in different words). While we did not fine-tune the model on a specific training dataset, frontier LLMs are pre-trained on millions of SQL queries. Inflating our dataset with simple queries would merely test if the model memorized basic SQL syntax during pre-training, rather than testing its ability to perform true compositional spatial reasoning.

As Finegan-Dollak et al. (2018) [6] demonstrated, redundancy inflates accuracy scores without actually testing compositional logic. They standardized several widely-cited single-domain benchmarks with as few as 128 queries (Yelp) and 131 queries (IMDB), proving that *query diversity and stratification matter more than raw volume*.

Instead of generating redundant variations of basic **SELECT** queries, our benchmark represents structurally distinct spatial challenges. Passing this compact, highly stratified test proves the model’s fundamental spatial reasoning capabilities much better than passing redundant variations of the same template. Furthermore, prior to GeoSQL-Eval [5], no standardized benchmark existed for evaluating LLMs on PostGIS spatial queries, making our focused dataset a novel contribution to an emerging research area.

4.3 The English SQL Benchmark (50 Queries)

We manually wrote 50 complex English questions to stress-test the AI’s ability to think geographically. We designed these questions across five difficulty tiers to see exactly where the AI would fail:

- **Easy:** Direct lookups (e.g., “Where is the library?”).
- **Medium:** Multiple filters (e.g., “Find cafes or restaurants”).
- **Hard:** Tricky negative logic (e.g., “Find a park but not a sports pitch”).

- **Extra (Intent):** Translating human feelings into map locations (e.g., “I am stressed, find a calm place”).
- **Extreme (Spatial Math):** Forcing the AI to use complex geographic math (e.g., “Find a cafe within 500m of the library”).

4.4 The Multi-Lingual Translation Dataset (56 Queries)

To test the translation models, we wrote 56 complex location requests in multiple Indian languages (**Hindi, Tamil, Telugu, Marathi, and Punjabi**). This dataset forced the AI to deal with messy regional dialects. We broke it into three categories:

- **Native Script:** Perfect, formal language written in native scripts (e.g., Devanagari).
- **Romanized (Hinglish):** Conversational internet text written with English characters (e.g., “Mujhe bhook lagi hai...”).
- **Colloquial Slang:** Highly localized phrases that have no direct English translation (e.g., “koi chai ki tapri hai kya?”).

We evaluated these queries as raw text first. Then, we generated synthetic audio for all of them to test how well the AI could recognize and translate spoken Indian languages.

5 Evaluation and Result(s) (Comparisons)

5.1 Text-to-SQL Model Selection

We evaluated four open-weight models on the 50-query Golden Benchmark. All models were evaluated locally using a standard prompt without model-specific “cheat codes” or examples, testing their zero-shot reasoning capabilities on the OSM schema.

Model	Pass Rate	Fail Rate	Avg Struct Score	Avg F1 Score
Qwen2.5-Coder-7B-Instruct	78.0% (39/50)	4.0% (2/50)	0.854	0.780
SQLCoder-8B (llama-3-sqlcoder)	46.0% (23/50)	54.0% (27/50)	0.600	0.459
CodeS-7B (BIRD-SQL Champ)	32.0% (16/50)	68.0% (34/50)	0.697	0.385
DeepSeek-Coder-7B	28.0% (14/50)	72.0% (36/50)	0.489	0.297

Table 1: Overall Text-to-SQL Performance Comparison

5.1.1 Conclusion and Root Cause Analysis

Qwen2.5-Coder-7B emerged as the definitive winner, displaying immense spatial reasoning capabilities despite its small parameter footprint.

- **SQLCoder-8B Failures:** SQLCoder suffered from fatal weaknesses. It mapped wrong columns without explicit schema hints (e.g., mapping “library” to `leisure` instead of `amenity`), struggled with negation (using `IS NULL` instead of `NOT ILIKE`), and completely broke down on spatial joins, writing invalid Common Table Expressions (CTEs).
- **CodeS/DeepSeek Failures:** CodeS-7B (the BIRD-SQL academic champion) completely failed (32% pass rate). Because it was aggressively fine-tuned on complex relational schemas, it over-complicated our single denormalized OSM schema and hallucinated non-existent columns (e.g., `address`).
- **Qwen2.5-Coder Success:** Qwen natively understood spatial semantics, generating clean spatial JOINS and `ORDER BY ST.Distance` logic without hallucinating.

Final Recommendation for Text-to-SQL: Qwen2.5-Coder-7B-Instruct is the recommended model for denormalized spatial logic.

5.2 Testing the Translation AI

Since our SQL AI only speaks English, we needed a translator for our Indian users. We tested several translation models on our 56-query multi-lingual dataset.

Model	Overall Pass Rate	Hinglish/Slang Pass Rate	Total SQL Failures
Google Translate (Cloud Baseline)	96.4%	71.4%	0
Sarvam-Translate (4B)	87.5%	71.4%	2
NLLB-200 (1.3B)	85.7%	57.1%	3
SeamlessM4T v2 (Large)	83.9%	42.9%	High

Table 2: Translation Model Performance across Dialects

5.2.1 Why Sarvam Won

While all models did well on perfect, dictionary-standard Hindi, they struggled with internet slang.

- **NLLB-200 & SeamlessM4T (Failed):** Both of these models completely failed on internet slang. They would often just repeat the Hindi slang word back in English letters (leaving words like “tapri” untranslated), which caused the SQL AI to crash.
- **Sarvam-Translate (Success!):** Sarvam proved to be the best text translator. It flawlessly translated messy slang into proper English. Its only flaw is that it is a bit “chatty” (sometimes it tries to hold a conversation instead of just translating), so we had to write a strict prompt to force it to behave.

Final Recommendation: Sarvam-Translate (4B) is our official choice for translating messy text.

5.3 Testing Voice Recognition (ASR)

To allow users to literally speak to the map, we tested four different voice recognition engines using our synthetic audio dataset.

Model	Pass Rate	Avg Word Error Rate (WER)	Type
Sarvam API (Saaras)	83.9% (47/56)	0.339	Closed API
SeamlessM4T v2 (Large)	82.1% (46/56)	0.470	Local / Open
Google Speech Recognition	78.6% (44/56)	0.514	Cloud API
OpenAI Whisper (Large)	66.1% (37/56)	0.632	Local / Open

Table 3: ASR Performance on Multilingual Spatial Audio

5.3.1 Why SeamlessM4T Won for Voice

- **OpenAI Whisper (Failed):** Despite being famous, Whisper was terrible at understanding regional Indian languages, completely failing on Telugu and Marathi.
- **Google Speech (Failed):** Google was decent, but completely broke down when a user mixed Hindi and English in the same sentence.

- **Sarvam API (Runner-Up):** Sarvam was incredibly accurate, but it requires an internet connection to use their API, which breaks our “100% offline” goal.
- **SeamlessM4T v2 (Success!):** Seamless was the massive winner. It listened to Indian audio and translated it directly into English text in a single step, entirely offline on our local GPU.

Final Recommendation: **SeamlessM4T v2** is our official choice for the voice engine because it is fast, accurate, and runs 100% offline.

5.4 The Final End-to-End App

The final product of this project is a fully integrated app that runs completely offline on a single local GPU. Here is how data flows through the system:

1. **User Input:** A simple web interface lets the user type a question or speak directly into their microphone in their native language.
2. **Voice-to-English (SeamlessM4T):** If the user speaks in Hindi, Marathi, or another language, the SeamlessM4T AI instantly transcribes the audio and translates it into English text entirely offline.
3. **Small-Talk Filter (Qwen):** Before doing any database math, the Qwen AI acts as a bouncer. If the user just says “Hello” or makes small talk, Qwen replies conversationally without wasting time searching the map database.
4. **Generating SQL (Qwen):** If the user asks a real spatial question (e.g., “Where is a quiet place?”), Qwen translates that English intent into strict PostGIS database code.
5. **Database Search:** The SQL code is executed on the local map database, returning the exact coordinates of the target locations.
6. **Reverse Translation:** Qwen looks at the database results and writes a friendly summary in English. That English summary is then passed back to SeamlessM4T, which translates it back into the user’s original language.
7. **Map & Voice Output:** Finally, the locations are plotted as interactive pins on a map. At the exact same time, SeamlessM4T generates a synthetic voice clip to speak the translated answer back to the user out loud.

6 Conclusion

We successfully built and tested an Intent-Driven Spatial Search Engine. By combining Qwen’s coding ability with SeamlessM4T’s translation skills, we proved that we can abandon rigid, exact-match text search in favor of intelligent, mathematical map routing.

Our custom grading framework proved that this architecture works. Our AI achieved a 78% success rate on highly complex geographic queries, completely outperforming models that were specifically designed for SQL. Most importantly, by integrating voice recognition and an interactive map, we created an app that lets humans search the world exactly how they think about it—conversationally and naturally.

6.1 Future Work

While our offline prototype is fully functional, here is how we can make it even better:

- **Make it Faster (Real-time Latency):** Currently, running two massive AI models on a single laptop GPU takes about 30 seconds to answer a question. To make the app feel like a real-time conversation, we need to speed this up. We can do this by using faster AI engines (like **vLLM**) or by streaming the audio so the AI starts speaking before it even finishes thinking.
- **Remember Past Questions (State Memory):** Right now, the AI forgets what you asked previously. If we give the AI memory, a user could ask “Find quiet places,” and then follow up with “Are any of them open right now?” The AI would remember the first search and just add a time filter, rather than starting entirely from scratch.
- **Use Live GPS:** Right now, the user has to say “near the library.” In the future, we can hook the app up to the phone’s live GPS so the AI automatically knows what “near me” means without asking.
- **Add Pictures (Street View):** Instead of just showing map pins, we could integrate Google Street View or generate images so the user can actually see what the location looks like before walking there.

References

- [1] Zhong, V., Xiong, C., & Socher, R. (2017). *Seq2SQL: Generating Structured Queries from Natural Language using Reinforcement Learning*. arXiv preprint arXiv:1709.00103.
- [2] Yu, T., Zhang, R., Yang, K., et al. (2018). *Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL Task*. Proceedings of EMNLP 2018. <https://aclanthology.org/D18-1425.pdf>
- [3] Li, J., Hui, B., Qu, G., et al. (2023). *Can LLM Already Serve as A Database Interface? A BIG Bench for Large-Scale Database Grounded Text-to-SQLs*. NeurIPS 2023. <https://bird-bench.github.io/>
- [4] Chang, S., Wang, J., Dong, M., et al. (2023). *Dr.Spider: A Diagnostic Evaluation Benchmark for Robust Text-to-SQL*. ICLR 2023. <https://openreview.net/forum?id=Wc5bmZo2v8>
- [5] Jiao, H., et al. (2026). *GeoSQL-Eval: First Evaluation of LLMs on PostGIS-Based NL2GeoSQL Queries*. Expert Systems with Applications. <https://haoyuejiao.github.io/GeoSQL-Eval-Leaderboard/>
- [6] Finegan-Dollak, C., Kummerfeld, J.K., Zhang, L., et al. (2018). *Improving Text-to-SQL Evaluation Methodology*. Proceedings of ACL 2018. <https://aclanthology.org/P18-1033/>
- [7] Defog AI. (2024). *SQLCoder: State-of-the-art LLM for Text-to-SQL Generation*. <https://defog.ai/sqlcoder/>
- [8] Alibaba Cloud. (2024). *Qwen2.5-Coder: A Large Language Model for Code Generation*. <https://github.com/QwenLM/Qwen2.5-Coder>
- [9] Sarvam AI. (2024). *Sarvam-Translate: Multilingual Translation for Indian Languages*. <https://huggingface.co/sarvamai/sarvam-translate>
- [10] The PostGIS Development Group. *PostGIS: Spatial and Geographic Objects for PostgreSQL*. <https://postgis.net/>
- [11] OpenStreetMap contributors. *OpenStreetMap*. <https://www.openstreetmap.org/>

Reproducible Code

- **VISRI Project Repository (GitHub):** *Intent-Driven Spatial Search Engine* (Contains all evaluation results, datasets, and end-to-end notebooks). <https://github.com/mayank-pillai-99/Intent-Driven-Spatial-Search-Engine>
- **Final End-to-End Pipeline:** Open in Google Colab
- **Text-to-SQL Evaluation Sandbox:** Open in Google Colab
- **Translation Evaluation Sandbox:** Open in Google Colab
- **ASR Evaluation Sandbox:** Open in Google Colab